

DETECTING FAKE JOB POSTINGS USING BIDIRECTIONAL LSTM

 Aravind Sasidharan Pillai

 International Research Journal of Modernization in Engineering Technology and Science (IRJMETS), Volume 05, Issue 03, March 2023

“ Core Achievement

This study introduces a deep learning framework using Bidirectional LSTM networks to effectively detect fake job advertisements by combining structured metadata with multi-field textual features.

≡ Abstract (Simplified)

Online job portals are increasingly targeted by scammers who publish fraudulent job listings to steal personal data or money from job seekers. To address this issue, the paper presents an automated detection system powered by a Bidirectional Long Short-Term Memory (Bi-LSTM) deep learning network. By combining structured information (such as salary ranges and locations) with multi-field text descriptions (such as job title, company profile, and requirements), the system analyzes text context in both forward and backward directions. The proposed model achieves an accuracy of 98.71% and an ROC AUC score of 0.91, significantly outperforming traditional ensemble models.

Practical Takeaways

Job boards and recruitment platforms can integrate this multi-branch Bi-LSTM architecture into their post submission pipelines to automatically flag fraudulent listings in real time, protecting job seekers from scams and enhancing platform integrity.

The Problem

The proliferation of fraudulent job advertisements on internet recruitment portals leads to financial losses, identity theft, and wasted effort for job seekers.

Previous Limitations

Existing solutions rely heavily on traditional machine learning algorithms (like Naïve Bayes or Random Forest) or basic deep learning approaches that fail to fully capture sequence dependencies and bidirectional context in textual descriptions.

Gap Addressed

The paper fills the gap by developing a holistic deep learning architecture that integrates both structured metadata and multi-attribute unstructured textual descriptions using Bidirectional LSTMs.

Core Idea

Process textual fields through separate parallel Bidirectional LSTM streams to capture forward and backward semantic context, then merge these context representations with dense numerical/categorical metadata for final binary classification.

Key Steps

- 1 Data preprocessing to clean raw text (removing punctuation, special characters, and HTML tags) and one-hot encode categorical variables.
- 2 Conversion of textual attributes into dense numeric vector embeddings using TextVectorization and Embedding layers.
- 3 Sequential feature extraction via five parallel Bi-LSTM branches corresponding to job title, company profile, job description, benefits, and requirements.
- 4 Feature fusion in a Merging Layer that concatenates outputs from all Bi-LSTM branches and dense numerical layers.
- 5 Classification using a dense hidden layer with ReLU activation followed by a single-neuron output layer with a Sigmoid activation function.


What's Novel?

Using a multi-branch architecture that independently processes multiple individual textual attributes through dedicated Bi-LSTM streams before fusing them with structured numerical data.

Architecture Overview

The system processes a job posting through two main pathways: structured metadata is processed through dense layers, while unstructured text attributes (title, company profile, description, benefits, requirements) are tokenized, embedded, and fed into five parallel Bi-LSTM branches. Outputs from all streams are concatenated in a Merging Layer, routed through a ReLU-activated dense layer, and evaluated by a Sigmoid layer to produce a probability score of fraudulence.

Pipeline & Control Flow



Core Components

Numeric/Categorical Branch

Receives and transforms structured metadata such as salary range, company size, and binary flags using dense neural layers.

Text Vectorization & Embedding Layer

Tokenizes raw text inputs and maps integer tokens into dense vector embeddings capturing semantic context.

Parallel Bi-LSTM Layers

Extracts contextual patterns from text sequences in both forward and backward directions for each text field independently.

Merging Layer

Concatenates feature vectors from the numeric branch and all five Bi-LSTM text branches into a unified vector representation.

Dense & Sigmoid Classification Layers

Applies non-linear transformations via ReLU activation and computes the final probability score of job fraudulence using a sigmoid function.

Performance Comparison

The proposed Bi-LSTM model significantly outperformed baseline ensemble models including LightGBM (98.18% accuracy, 0.83 ROC AUC), Random Forest (97.04% accuracy, 0.70 ROC AUC), and Gradient Boosting (96.67% accuracy, 0.68 ROC AUC).

Datasets Used

- Kaggle Real / Fake Job Posting Prediction dataset (17,880 listings: 16,244 genuine, 1,636 fake)

Key Findings

- The Bi-LSTM model achieved an overall accuracy of 98.71% and an ROC AUC score of 0.91.
- The proposed model achieved a precision of 0.89, recall of 0.83, and an F1-score of 0.86 on the test dataset.
- The model reached early stopping criterion after just 6 training epochs.

Strengths

- Multi-input architecture effectively handles both structured categorical data and unstructured, multi-field text descriptions.
- Bidirectional memory layers effectively capture rich contextual representations from past and future sequence tokens.
- Demonstrates high recall and F1-score compared to traditional tree-based classification techniques.

Limitations

- Evaluated on a dataset with severe class imbalance (~9.1% fake listings), which can introduce bias toward the majority class.
- Deep learning models suffer from a lack of interpretability and explainability regarding classification decisions.
- Generalization to external job portals, different languages, or shifting domain distributions remains unverified.

I have completely parsed and summarized this paper. Feel free to ask me anything about its parameters, equations, logic, or findings. How can I assist you?

how can i enhance the gap of this paper and make it better

To enhance the research presented in this paper and make it a stronger, more robust study, you can directly address the limitations and future research directions that the author himself points out in **Section IV (Results and Discussion)**, specifically under **Section 4.3 (Discussion)**.

Here is a step-by-step guide on how you can bridge the gaps in this paper and elevate the research to the next level:

1. Advanced Handling of Class Imbalance

The Gap in the Paper: In **Section 2.2 (Description of Data)**, the author notes that the dataset is highly imbalanced, containing 16,244 genuine job postings but only 1,636 fake ones (about a 10:1 ratio). In **Section 4.3 (Class Imbalance)**, the author admits this imbalance can lead to biased models and suggests oversampling or undersampling for future work.

How to Enhance It:

Implement SMOTE or ADASYN: Use Synthetic Minority Over-sampling Technique (SMOTE) to generate synthetic examples of the minority class (fake jobs) before training.

Apply Class Weights / Focal Loss: Modify the loss function of the Bi-LSTM model to penalize misclassifications of the "fake" class more heavily.

Use Data Augmentation: Use LLMs (like GPT-4) or back-translation (translating English text to French and back to English) to generate realistic, paraphrased fake job descriptions to balance the dataset naturally.

2. Upgrade the NLP Architecture (Transition to Transformers)

The Gap in the Paper: The paper relies on a **Bidirectional LSTM** (discussed in **Section 3.2**). While Bi-LSTMs are great for sequential data, they have largely been superseded in modern NLP by Transformer-based models.

How to Enhance It:

Compare with BERT/Roberta: Train a pre-trained transformer model like BERT, RoBERTa, or DeBERTa on this dataset. Transformers use self-attention mechanisms that capture long-range text dependencies much better than LSTMs.

Hybrid Models: Combine Transformer embeddings (like BERT embeddings) with numerical/categorical metadata (using the merging layer concept from **Figure 4**) to see if it outperforms the paper's Bi-LSTM + TextVectorization approach.

3. Incorporate Explainable AI (XAI)

The Gap in the Paper: In **Section 4.3 (Interpretability and Explainability)**, the author notes that deep learning models are "black boxes" and lack interpretability, which is a major issue when dealing with sensitive hiring data.

How to Enhance It:

Apply LIME or SHAP: Integrate LIME (Local Interpretable Model-agnostic Explanations) or SHAP (SHapley Additive exPlanations) into your model. This will allow you to highlight *exactly* which words (e.g., "wire transfer," "no experience needed," "urgent") or features (e.g., missing

company logo) contributed most to the model classifying a job as "fake."

4. Test Cross-Platform Domain Adaptation

* **The Gap in the Paper:** In **Section 4.3 (Transfer Learning and Domain Adaptation)**, the author states that the model was only trained and tested on one specific Kaggle dataset, leaving its performance on other job boards or languages uncertain.

* **How to Enhance It:**

* **Multi-source Evaluation:** Scrape a new, completely independent dataset of job postings from LinkedIn, Indeed, or Glassdoor in 2024.

* **Zero-Shot / Few-Shot Testing:** Test how well the model trained on the Kaggle dataset performs on this new real-world dataset without retraining. This will prove whether your model actually generalizes well or if it just memorized the Kaggle dataset.

5. Deeper Feature Engineering on Metadata

* **The Gap in the Paper:** While the paper uses categorical data like "location" and "industry" (using one-hot encoding as described in **Section 2.4**), it treats them very simply.

* **How to Enhance It:**

* **URL/Domain Verification:** If a job posting includes a contact email, write a script to check if the email domain matches the official company website (e.g., a recruitment email from a Gmail account vs. an official `@company.com` account is a massive red flag).

* **Geographical Anomaly Detection:** Check if the location listed in the job posting matches the actual location of the company's headquarters.

Summary of your potential research paper:

If you implement these steps, your paper could be titled: **"An Explainable Transformer-Based Approach to Fake Job Detection with Cross-Domain Validation."** You would be directly solving the black-box issue, upgrading the tech stack from LSTM to Transformers, and proving your model works in the real world on fresh data.